

Convolutional Neural Nets

Tuesday, November 26, 2019 1:07 PM

Extra reading:

<https://pdfs.semanticscholar.org/450c/a19932fcef1ca6d0442cbf52fec38fb9d1e5.pdf>

http://ais.uni-bonn.de/papers/icann2010_maxpool.pdf

Steps:

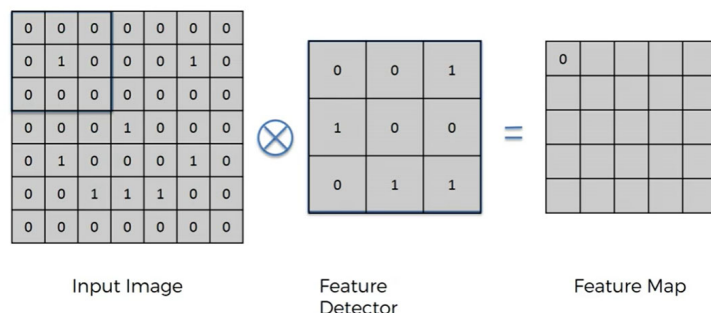
1. Convolutional operation
2. ReLU layer
3. Pooling
4. Flattening
5. Full ocnnnection

A convolution is a combined integration - shows how one function modifies the shape of another

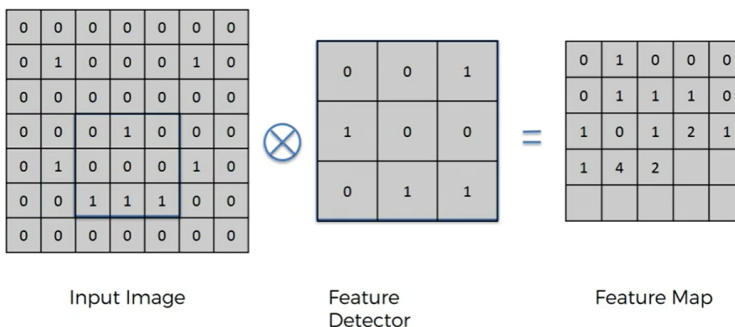

$$(f * g)(t) \stackrel{\text{def}}{=} \int_{-\infty}^{\infty} f(\tau) g(t - \tau) d\tau$$


Feature detector:

- Usually 3x3 matrix
- "Filter"
- Denoted




Each movement of the filter along the input image is called a stride and maps to some output called a "Feature map" or "Convolved Feature" or "Activation map"





Step 1b: Rectify to break up linearity




Images themselves are highly non-linear.

- Removes linear gradients in color
- Ex. Replaces white to black gradients with just white

Step 2: pooling

Adds spatial invariance

- Accounts for distorted features
- Removing information prevents overfitting




Max pooling strides across each feature map and takes max value

Subsampling takes the average across each stride

Step 3: Flattening

Transform pooled map into a vector




Flattening




Step 4: Full Connection

All layers are connected and Final layer votes on output



Softmax function:



$$f_j(z) = \frac{e^{z_j}}{\sum_k e^{z_k}}$$

Softmax (normalized information function) function normalize output so that output confidence adds up to 1


Prevents values like 45% cat and 90% dog

Cross entropy function:

$$L_i = -\log\left(\frac{e^{f_{y_i}}}{\sum_j e^{f_j}}\right)$$

$$H(p, q) = -\sum_x p(x) \log q(x)$$

- Alternative to max squared sum function
- "Loss" function... not a "cost" function like sq sum
- Minimize loss function maximize prediction performance



Dog 0.9
Cat 0.1

$$H(p, q) = -\sum_x p(x) \log q(x)$$



